

# Local RAG Pipeline with Kubernetes & Ollama

A fully local, containerized Retrieval-Augmented Generation (RAG) system built from first principles using Python, Docker, and Kubernetes. This project avoids high-level RAG abstractions in favor of explicit, inspectable system boundaries.

## 1. Overview

This repository implements a local-first RAG pipeline that ingests documents, builds a persistent index, retrieves relevant context, and generates grounded answers using a local LLM. All components run locally with no cloud dependencies.

## 2. System Architecture

### High-Level Flow

```
PDF Docs
  |
  Seed Docs Job
  |
  Docs PVC
  |
  Ingest Job (chunk + embed)
  |
  Index PVC (local_index.pkl)
  |
  Query Job
  |
  ollama LLM Deployment
```

### Component Responsibilities

seed-docs	: Copy PDFs into persistent storage
rag-ingest	: Chunk documents and build index
rag-query	: Retrieve context and generate answer
ollama	: Local LLM inference service
PVCs	: Persistent documents, index, and model weights

## 3. Model & Algorithm Choices

Chunking is deterministic and transparent. Embeddings are computed locally and stored in a serialized index. Retrieval uses top-K similarity search with full visibility into scores and source metadata.

### LLM Backend

Ollama is used for fully local inference. The llama3.1:8b model was selected to balance reasoning capability with local resource constraints.

## 4. Why Docker and Kubernetes

Docker ensures reproducible execution environments and prevents silent dependency drift. Kubernetes cleanly separates batch workloads (ingest/query) from long-lived services (LLM inference), while making storage dependencies explicit through PersistentVolumeClaims.

## 5. Execution Flow

```
kubectl apply -f k8s/00-namespace.yaml  
kubectl apply -f k8s/01-pvc.yaml  
kubectl apply -f k8s/02-ollama.yaml  
kubectl apply -f k8s/03-seed-docs-job.yaml  
kubectl apply -f k8s/04-job-ingest.yaml  
kubectl apply -f k8s/05-job-query.yaml
```

## 6. Repository Structure

```
llm_project_rag_local/  
    data/  
    docs/  
    rag_local/  
    scripts/  
    k8s/  
    Dockerfile  
    requirements.txt  
    README.md
```

## 7. Project Rationale

This project demonstrates systems-level understanding of RAG pipelines: explicit data flow, persistence boundaries, failure modes, and separation of concerns. It favors engineering clarity over abstraction or novelty.